

Balmung: Access and pricing

Why are prices so low on Balmung?

Sometimes, I play the video game Final Fantasy XIV with friends. Final Fantasy is an MMO with a fairly active economy selling in-game items for in-game currency, the *marketboard*.

Players of Final Fantasy XIV are split into servers, each with its own marketboard. You can only *sell* on the marketboard of the server you are based on, but usually *buying* on the marketboard of another server is only a minor inconvenience. So there's usually not much difference in pricing from server-to-server: players would rather spend a little bit of time traveling to another server rather than eating a hefty convenience fee.

But there's one major exception: the Balmung server. Balmung is so active and so over-populated that travel to it is almost always disabled. So players based on Balmung can buy anywhere, but sellers on Balmung have a much more restricted audience of only Balmung-based players.

And prices on Balmung are often *much* lower than on any other server.

Intuitively, it seems like these facts should be related: a smaller pool of potential buyers artificially reduces demand, therefore reduces prices. But what is the actual mechanism that drives prices down on Balmung? Certainly players aren't doing supply/demand calculations in their head when they set prices for the virtual items they are selling.

This post is an attempt to investigate.

i tl;dr! Click me to see a solution summary

Prices everywhere move with a random walk. If the price on Balmung is higher than on other servers, buyers can simply leave Balmung and buy elsewhere. But if the price on Balmung is lower than elsewhere, sellers have no recourse. So prices on Balmung move in a one-sided random walk that is always below the equilibrium price.

Is it one-sided price drift?

No item sold on the marketboard is priced using sophisticated theory. Most items, it seems, are priced using the following algorithm:

1. Look at the currently listed prices for the item you're selling
2. Price your item one gil (one dollar) less than the cheapest competitor.

Let $b(P)$ be the chance that a player wants to buy our item at price P , and $s(P)$ the chance that they want to sell our item at price P .

The equilibrium price, P_{eq} , has to satisfy

$$s(P_{eq}) = b(P_{eq})$$

If demand and supply are both high, the equilibrium price should be reached with minimal randomness:

$$P_{actual} = P_{eq} + \epsilon$$

But what happens if demand is not that high? Then the actual price should drift around the equilibrium price appreciably as, by random chance, there is higher supply one day and higher demand the next:

$$P_{actual} = P_{eq} + \text{drift}$$

OK, so what happens if “the rest of the servers” have high demand but Balmung alone has low demand? You should expect one-sided drift in the price.

If there are ever too many Balmung *buyers*, $P_{Balmung} > P_{elsewhere}$, buyers can simply leave Balmung and buy elsewhere.

But Balmung sellers have no such recourse. If $P_{Balmung} < P_{elsewhere}$, they just have to accept $P_{Balmung}$.

Therefore, there is a sort of one-sided random walk occurring in Balmung: prices can get up to P_{eq} , but never go above them.

Hypothesis. Prices are lower on Balmung because of one-sided drift.

Modeling the effect of drift

Let's assume that we never go too far from equilibrium, so the arrival rates of sellers and buyers are both the equilibrium rate.

So our basic model for the drift is that buyers and sellers arrive with a rate R_{eq} *unless* $P_{Balmung} \geq P_{eq}$, in which case buyers arrive with rate 0 and sellers arrive with the same rate, R_{eq} .

This process is very close to a random walk. Because we have set the sell/buy rates to be equal, this random variable is actually $|X|$, where X is a random variable representing a random walk. So for larger times T , drift is described by a *folded normal distribution*, $\frac{Drift_T}{\sqrt{T}} \sim FoldN(0, \sigma)$.

The mean of the folded normal distribution is nonzero:

$$\mathbb{E}[FoldN(0, \sigma)] = \sigma \sqrt{\frac{2}{\pi}}$$

So in our simple model, average drift should grow unbounded with time,

$$\mathbb{E}[Drift_T] \sim \sigma \sqrt{\frac{2T}{\pi}}$$

Which means that this simple model can't be valid for large enough times.

Adding seller reluctance

As price goes lower, the amount of sellers will reduce. Since Final Fantasy players are just having fun, they're far from rational market agents and so the exact way in which their selling reduces will likely be very complicated.

Let's choose a very simple model, where the rate of selling at price P is just

$$R_P = R_{eq} \times \frac{P}{P_{eq}}$$

this has the pleasant features of being 0 at price 0 and R_{eq} at $P = P_{eq}$.

At this point, our model is complicated enough that we should just write it in code. A closed-form solution exists for our choice of R_P , but after all we might want to choose another R_P and so simulating seems like an efficient way forward.

```

function walled_biased_walk(time, eqprice, stepsize=1)
    function step_fn(w_prev, _)
        r = rand() < w_prev/(eqprice + w_prev) ? -1 : 1
        w_new = min(max(w_prev + r*stepsize, 0), eqprice)
        return w_new
    end
    walk = accumulate(step_fn, 1:time-1; init=eqprice)
    return walk
end

using Plots
plot(walled_biased_walk(10000,200),title = "Example walk", label = "Price")

```



Figure 1: Example of a walled biased walk with $P_{eq} = 200$ and $T = 10000$

Now we just need to calculate the expected value of this process.

```

using Statistics
function walled_biased_walk_mean(time, eqprice, stepsize=1)
    W = walled_biased_walk(time, eqprice, stepsize)
    return mean(W)
end

```

```
means = [200 - walled_biased_walk_mean(10000, 200) for _ in 1:1000] # initialize with one walk
using Plots
plot(means, label="Mean of 1000 walks", xlabel="Walk number", ylabel="Mean price", title="Mean price discount , P = 200")
```

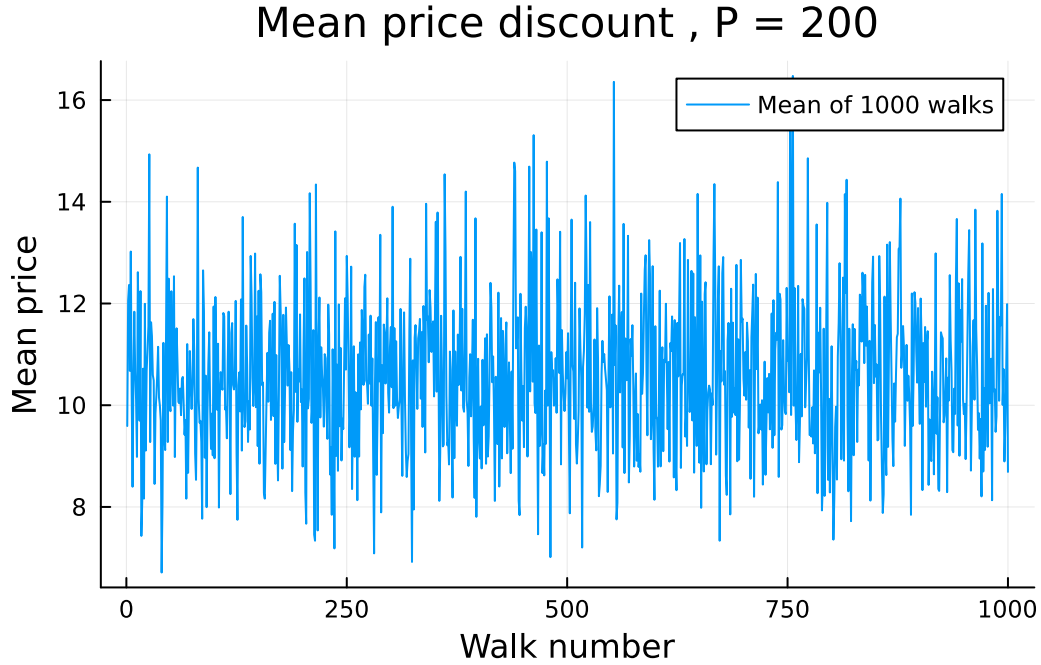


Figure 2: Mean discount for the biased walk with $P_{eq} = 200$ and $T = 10000$

Model parameters

What choices can we make in this model? Given our assumptions, there are two real parameters left in the model:

1. The choice of R_P
2. The step size s , also possibly a function of P , $s(P)$.

Both these parameters can be quite general functions (under some constraints, e.g. $R_P(0) = 0$ and $R_P(P_{eq}) = R_{eq}$).

But let's make an assumption of linearity, assuming $R_P = R_{eq} \frac{P}{P_{eq}}$ and $s(P) = \alpha P_{eq} + \beta$ for some constants α and β .

A constant step size ($\alpha = 0, \beta \neq 0$) will lead to a constant discount despite the size of P , which seems unrealistic: the amount of drift should scale with the price of the item – sellers

are willing to drop the price 100 gil (dollars) on a 100 000 gil item if they're worried about selling it, but not if it's a 100 gil item.

So we must have $\alpha > 0$.

So to start with basic modeling, let's choose a step size that is a fraction of P_{eq} , say $s = \alpha P_{eq}$ and think about β later.

```
using Plots
using Statistics
function avg_discount(time, eqprice, stepsize = 1, nwalks = 1000)
    means = [eqprice - walled_biased_walk_mean(time, eqprice, stepsize) for _ in 1:nwalks]
    return mean(means)/eqprice
end
discounts = Float64[]
for p in 50:50:5000
    push!(discounts, avg_discount(1000, p, 0.02*p))
end
plot(50:50:5000, discounts,
     title = "Discount/P for fixed step size 0.02P",
     xlabel = "Equilibrium price P",
     label = "Mean discount over P",
     ylabel = "Mean discount over P",
     ,ylims = (0, 0.3))
```

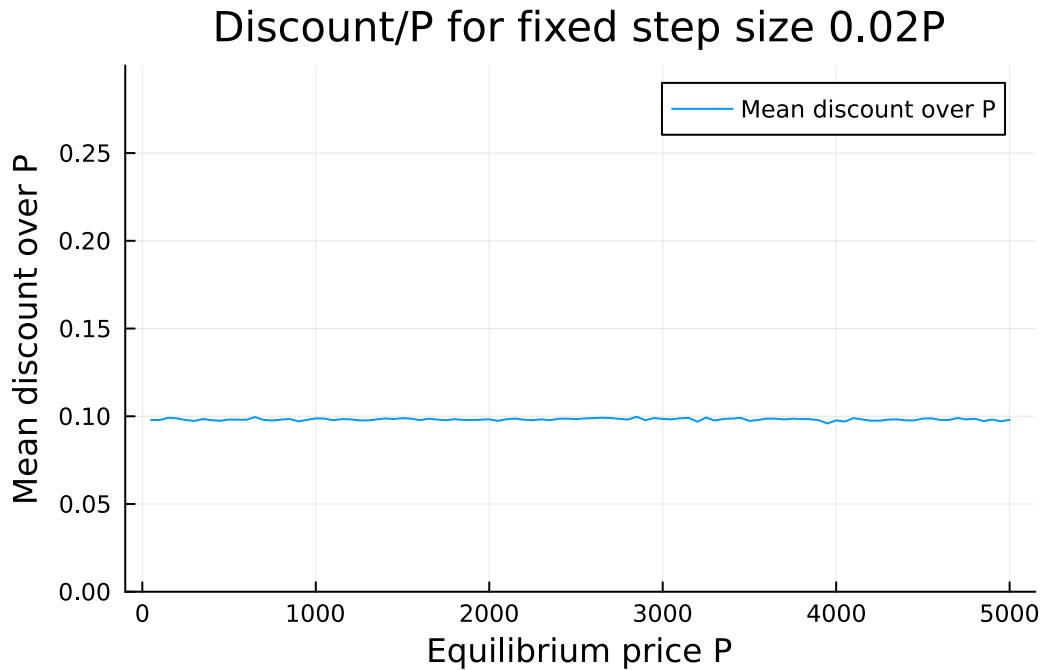


Figure 3: Mean discount/ P_{eq} is relatively constant no matter price, for fixed step size $0.02P_{eq}$

It's very motivating to see that the mean fractional discount in our model is relatively constant over a wide range of prices. Now, what about if we vary the step size?

```
using Plots
using Statistics
function avg_discount(time, eqprice, stepsize = 1, nwalks = 1000)
    means = [eqprice - walled_biased_walk_mean(time, eqprice, stepsize) for _ in 1:nwalks]
    return mean(means)/eqprice
end
discounts = Float64[]
for s in 0.01:0.01:0.25
    push!(discounts, avg_discount(1000, 1000, s*1000))
end
plot(0.01:0.01:0.25, discounts,
     ylims = (0, 0.3),
     xlabel = "Step size coefficient",
     ylabel = "Mean discount over P",
     label = "Mean discount/P",
     title = "Mean price discount/P with P = 1000")
```

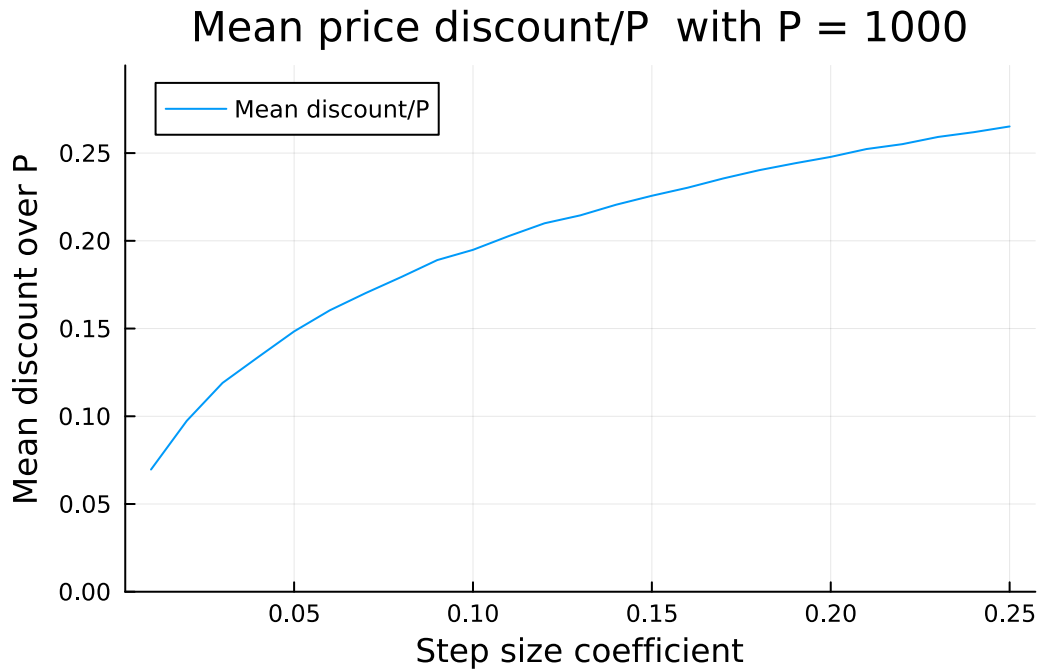


Figure 4: Discount/ P_{eq} vs step size coefficient for walled biased walk with $P = 1000$

The dependence seems roughly proportional to $\sqrt{\text{step size}}$:

Finally, let's look at the effect of β for fixed $\alpha = 0.1$:

```
discounts = Float64[]
for beta in 0:1:100
    push!(discounts, avg_discount(1000, 1000, 0.1*1000 + beta))
end
plot(0:1:100, discounts,
     ylims = (0,0.3),
     xlabel = "Beta coefficient",
     ylabel = "Mean discount over P",
     label = "Mean discount/P",
     title = "Mean price discount/P with P = 1000 and alpha = 0.1")
```


Above, but with sqrt(step size)

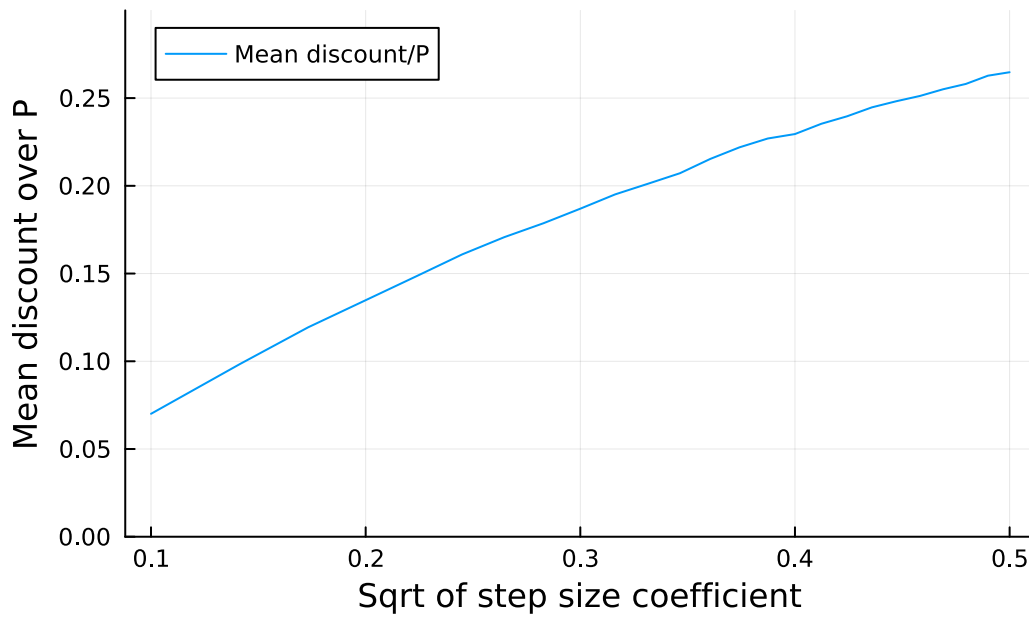


Figure 5: Discount/ P_{eq} is roughly linear in sqrt(step size)

Mean price discount/P with $P = 1000$ and alpha =

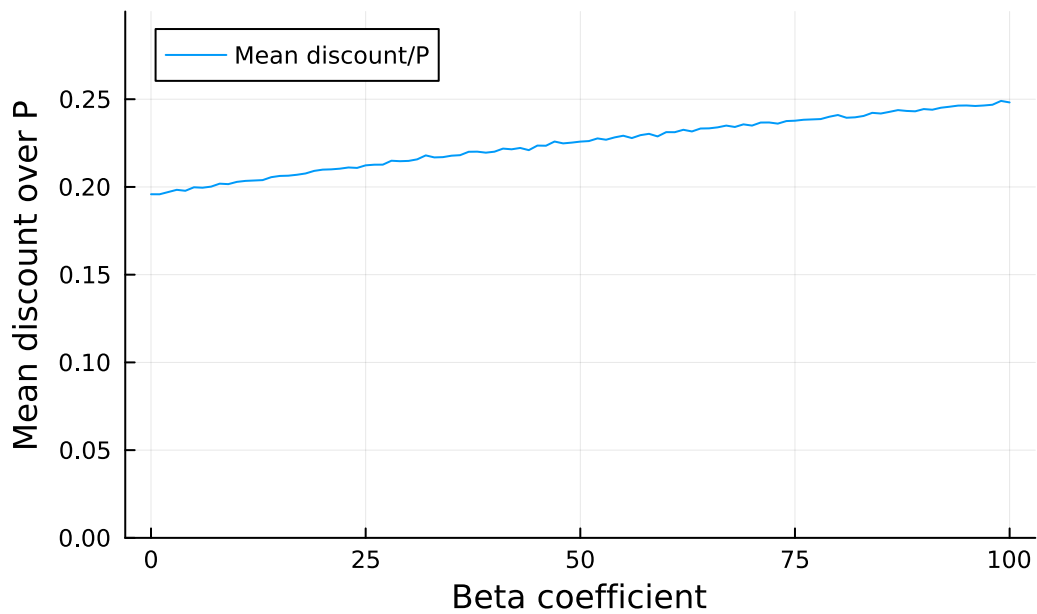


Figure 6: Discount/ P_{eq} vs beta

So β has a small, linear effect on mean discount.

Does my model match the data?

Not every commodity will have the ‘Balmung discount’ effect as I’ve modelled it, because the model as presented above is pretty simplistic. For example, it assumes the price on every other server remains exactly at equilibrium, which is only a reasonable approximation for highly traded items. While a Balmung discount might exist for items traded less frequently, modeling it would require adding in the assumption that P_{eq} itself can do a random walk (a fun problem, and possibly a future post!)

To start, let’s look at the data for a bucket of randomly chosen commodities.

```
using HTTP
using JSON
using DataFrames
ids = [44106,44347,13114, 13115, 47915,41763,41771]
discounts = Float64[]
worldV = Float64[]
dcV = Float64[]
for id in ids
    url = "https://universalis.app/api/v2/aggregated/Balmung/$id"
    response = HTTP.get(url)
    data = JSON.parse(String(response.body))
    avg_sale_price_world = data["results"][1]["nq"]["averageSalePrice"]["world"]["price"]
    avg_sale_price_dc = data["results"][1]["nq"]["averageSalePrice"]["dc"]["price"]
    daily_velocity_world = data["results"][1]["nq"]["dailySaleVelocity"]["world"]["quantity"]
    daily_velocity_dc = data["results"][1]["nq"]["dailySaleVelocity"]["dc"]["quantity"]
    balmung_discount = (avg_sale_price_dc - avg_sale_price_world) / avg_sale_price_world
    push!(discounts, balmung_discount)
    push!(worldV, daily_velocity_world)
    push!(dcV, daily_velocity_dc)
end
df = DataFrame(ID=ids, Discount=discounts, WorldVelocity=worldV, DCVelocity=dcV)
println(df)
```

7×4 DataFrame

Row	ID	Discount	WorldVelocity	DCVelocity
	Int64	Float64	Float64	Float64
1	44106	0.287494	1330.7	13622.9
2	44347	-0.471662	0.313327	296.721

3	13114	0.0322311	40.4192	434.898
4	13115	0.165654	61.7254	319.907
5	47915	-0.00829986	1.25331	7.83317
6	41763	0.522656	230.922	3367.63
7	41771	0.255139	220.582	1837.34

The above table shows the discount for each commodity, as well as the average number of daily units sold on Balmung (WorldVelocity) and the average number of daily units sold on all other servers (DCVelocity).

This bucket is very encouraging. All three highly sold commodities have a consistent discount of about 25. It's commodities like these that we claimed to model anyhow. (With a more complicated model, we might expect to be able to model the smaller discount of the moderately sold commodity 13115, but the very rarely sold commodities are likely not governed by a random walk at all.)

Appendix: A predictive model, and the perils of limited data

OK, so I purport to have a model explaining the Balmung discount, but the real test of any model is whether it can predict unknown data. Our simple model has some parameters (that a more principled model might be able to partially predict, but for now we will just learn them from the data). We'd like to learn them by fitting our model to some data, then use the learned parameters to predict the discount for a new commodity.

One issue is that the Final Fantasy market is not *that* large; there are not so many commodities traded in sufficient volume to help us fit our model. Further, the API to access marketboard data is heavily rate-limited, so collecting large amounts of data is time-consuming. So let's restrict ourselves to learning limited parameters and hoping our loss is not too large.

So let's restrict ourselves to learning limited parameters with limited data. There are three things we could learn from the data that seem useful:

1. The step size $step(P_{eq}, wV, dcV)$, which is a function of the equilibrium price and the velocities on both servers.
2. The minimum (wV, dcV) at which our model applies.
3. The function $R_P(P, P_{eq}, wV, dcV)$.

The function R_P , I claim, is commodity-specific. The reason is that commodities are produced for different reasons. For a real life example, think about chicken wings. Chickens are raised primarily for breast meat, and wings are a byproduct. So, for chicken wings, supply shouldn't vary much with the price of wings: it mostly depends on the price of breast meat.

Therefore, learning a 'universal' R_P won't do us much good, and so our simple linear model seems like a reasonably good choice at our current level of sophistication.

So we will instead focus on 1), 2). Again, let's choose something very harsh and set an arbitrary binary threshold for velocities and just learn 1), leaving dealing with 2) to a future post where we deal with a more sophisticated model which incorporates the velocities directly. Then we just need to learn 1).

Collecting training data

Collecting data is a little tricky, since the API is *heavily* rate-limited and doesn't natively provide a way to get commodities by velocity.

I tried uniformly sampling market data, but the API is very slow and most of the items I sampled had too low velocity to be modeled. Instead, we sample from the most recently updated items on the marketboard.

```
using HTTP, JSON, StatsBase

# uniformly sample market data
function collect_market_data(item_ids, batch_size=100, num_samples=1000, world="Balmung", min_velocity=0.0, min_price=0.0)
    # Sample item IDs
    sample_ids = StatsBase.sample(item_ids, num_samples, replace=false)
    # Collect columns as vectors
    cols = Vector{Vector{Float64}}{0}

    # Batch request
    for batch in Iterators.partition(sample_ids, batch_size)
        id_str = join(batch, ",")
        url = "https://universalis.app/api/v2/aggregated/$world/$id_str"
        response = HTTP.get(url)
        data = JSON.parse(String(response.body))
        # The response has a "results" array
        for item in data["results"]
            try
                avg_sale_price_world = item["nq"]["averageSalePrice"]["world"]["price"]
                avg_sale_price_dc = item["nq"]["averageSalePrice"]["dc"]["price"]
                daily_velocity_world = item["nq"]["dailySaleVelocity"]["world"]["quantity"]
                daily_velocity_dc = item["nq"]["dailySaleVelocity"]["dc"]["quantity"]
                # if the velocities are too low, skip this item
                # if the price is too low, skip this item
                if (daily_velocity_world < min_velocity) || (avg_sale_price_dc < min_price)
                    continue
                end
                push!(cols, [avg_sale_price_world, avg_sale_price_dc, daily_velocity_world, daily_velocity_dc])
            catch
                continue
            end
        end
    end
end
```

```

        catch
            continue # skip items with any missing data
        end
    end
end

# Convert collected columns to a matrix
return hcat(cols...) # 4 x N matrix
end

# sample from the most recent market data
worldList = ["Balmung", "Mateus", "Gilgamesh", "Zalera", "Exodus", "Sargatanas", "Jenova", "C"]
function collect_recent_market_data(min_velocity=100, entries=200, minprice = 300, worlds = worldList)
    all_items = []
    for world in worlds
        url = "https://universalis.app/api/v2/extra/stats/most-recently-updated?world=$world"
        response = HTTP.get(url)
        item_ids = map(item -> item["itemID"], JSON.parse(String(response.body))["items"])
        append!(all_items, item_ids)
    end
    num_ids = length(all_items)
    return collect_market_data(all_items, 100, min(num_ids, 1000), "Balmung", min_velocity, minprice)
end

```

collect_recent_market_data (generic function with 5 methods)

Modeling the step size

Our model predicts a discount, so our loss function will be the mean squared error between the predicted fractional discount and the observed fractional discount.

Now we need to fit the step size function. One approach would be to use a simple neural network to incorporate the price and velocity as features.

First, let's collect some training data from the market and save it.

```

using CSV
using DataFrames
url = "https://universalis.app/api/v2/marketable"
response = HTTP.get(url)
item_ids = JSON.parse(String(response.body))
if !isfile("data.csv")

```

```

println("Collecting training data...")
data_matrix = collect_recent_market_data(600,200,400,worldList)
CSV.write("data.csv", DataFrame(data_matrix,:auto))
else
println("Training data already exists, skipping collection.")
data_matrix = Matrix(CSV.read("data.csv", DataFrame))
end

```

Training data already exists, skipping collection.

```

4×3 Matrix{Float64}:
 395.853  9038.65  1114.93
 460.044  9572.72  1119.71
 622.748   615.168   761.288
 5508.78  3659.45  7620.46

```

Train a neural network:

```

using Flux, HTTP, JSON, CSV

# Load the training data
data_matrix = CSV.read("data.csv", DataFrame)
# Convert to a matrix for processing
data_matrix = Matrix(data_matrix)

# Define a simple neural network for step size prediction
function step_size_model()
    return Chain(
        Dense(3, 4, relu), # Input layer: 3 features (dcPrice, worldVelocity, dcVelocity)
        Dense(4, 4, relu), # Hidden layer
        Dense(4, 1)        # Output layer: step size
    )
end

# Create the model
function loss(model, X, Y)
    preds = [avg_discount(1000,X[2, i], only(model(X[:, i])),20) for i in 1:size(X, 2)]
    return sum((preds .- Y).^2)
end

model = step_size_model()
# Define the optimizer
optimizer = Flux.setup(Adam(), model)

```

```

# Prepare the training data
X_train = data_matrix[2:4, :] # Features: dcPrice, worldVelocity, dcVelocity
Y_train = (data_matrix[1, :] .- data_matrix[2, :])./ (data_matrix[1, :]) # Target: fractional
# Reshape for Flux
X_train = reshape(X_train, size(X_train, 1), size(X_train, 2))
Y_train = reshape(Y_train, size(Y_train, 1), size(Y_train, 2))
# Training loop
epochs = 600
for epoch in 1:epochs
    Flux.train!(loss, model, [(X_train, Y_train)], optimizer)
    if epoch % 100 == 0
        println("Epoch $epoch: Loss = $(loss(model, X_train, Y_train))")
    end
end
end

```

```

Warning: Layer with Float32 parameters got Float64 input.
The input will be converted, but any earlier layers may be very slow.
layer = Dense(3 => 4, relu) # 16 parameters
summary(x) = "3-element Vector{Float64}"
@ Flux ~/.julia/packages/Flux/9PibT/src/layers/stateless.jl:60

```

```

Epoch 100: Loss = 0.1818477567518088
Epoch 200: Loss = 0.1790142582198929
Epoch 300: Loss = 0.18204088714897143
Epoch 400: Loss = 0.18077954746157335
Epoch 500: Loss = 0.17975902915370653
Epoch 600: Loss = 0.17987698314056505

```

So how good is our model?

Let's look at the results.

```

using DataFrames
# make a table of prediced vs actual discounts
X = X_train
preds = [avg_discount(1000, X[2, i], only(model(X[:, i])), 100) for i in 1:size(X, 2)]
actual_discounts = (data_matrix[2, :] .- data_matrix[1, :])./data_matrix[2, :] # actual fractional
df_results = DataFrame(
    BalmungPrice = data_matrix[1, :],
    RestPrice = data_matrix[2, :],
    PredictedDiscount = preds,
)

```

```

ActualDiscount = actual_discounts,
DiscountError = preds .- actual_discounts,
StepSize = [only(model(X[:, i])) for i in 1:size(X, 2)],
WorldVelocity = data_matrix[3, :],
DCVelocity = data_matrix[4, :]
)

```

	BalmungPrice	RestPrice	PredictedDiscount	ActualDiscount	DiscountError	StepSize	
	Float64	Float64	Float64	Float64	Float64	Float32	
1	395.853	460.044	3.87379e-6	0.139533	-0.139529	0.000149943	...
2	9038.65	9572.72	0.333534	0.0557899	0.277744	1040.35	...
3	1114.93	1119.71	3.13481e-6	0.00426952	-0.00426639	0.000149943	...

So the neural network approach didn't work very well. The core issue is that it turns out there are *very few* commodities which have sufficient velocity to be modeled. While I allowed the data to include intermediate-velocity examples, it's clear our model doesn't work for them. While those very few high-velocity commodities are the ones you keep seeing as a player, they are actually very sparse.

For those, we have few enough examples that the empirical table above shows $\alpha \approx 0.25$. We could leave satisfied with that, or we could try to develop a more sophisticated model that can handle intermediate-velocity commodities. As a consequence, we would have enough data to fit parameters numerically.